Consistent topics over time inside tweets

Data Mining Course Project Report

Davide Piva

2nd year CS Master Degree @ University of Trento

davide.piva-1@studenti.unitn.it

1 INTRODUCTION AND MOTIVATION

Twitter, as any other social network and blog, is a constant source of unstructured data that, if processed in the right way, can be leveraged to obtain valuable insights in several elds. Due to this reason, lots of companies have started to collect this huge amount of data in order to perform in-depth analysis on it; many of them have initiated to perform sentiment analysis to check customers satisfaction with respect to their products. Other organizations have started to use this asset, hereafter called Big Data, in their decision-making process that includes reporting, exploration of data and exploratory search (e.g., nding correlations). According to upGrad [18], Big Data can help create pioneering breakthroughs for companies that know how to use it correctly.

A challenging aspect in Big Data analysis is to extract valuable insights that can drive business strategies and decisions. In this context, the social network Twitter plays a crucial role during the information retrieval phase: millions of people publish everyday short messages, hereafter tweets, that often contain their desires or demands. The aim of many companies is to leverage these tweets in order to catch common issues or demands to guide their business strategies for their products or services. These insights, according to upGrad [18], can be used to develop new products/services, enhance marketing techniques, optimize customer service, improve employee productivity and nd radical ways to expand brand outreach.

In addition, organizations aim also to intercept occasional issues or demands that are not constant over time but are very popular in a bounded period of time. This is the case when there is a sudden issue or request for a particular product or service that a ects many customers. Catching this kind of insights, permits companies to better address their business short-term decision-making phase and enhance promptly their products or services. According to Ben Ridler [4], a business owner ability to e ectively deal with customer complaints provides a great opportunity to turn dissatis ed customers into active promoters of the business.

Nowadays, according to [17], there are more than 500 millions new tweets each day and this statistic is still growing every year. Thanks to this fact, many people have started to collect and group those tweets into datasets; many of them are easily and freely accessible by anyone through internet. Despite the fact that many of those datasets contain noisy data that force people to pre-process them, companies can elaborate those tweets in order to achieve the above mentioned business goals.

The aim of this project is to address in an e cient and e ective manner the procedure of nding consistent topics over time inside tweet texts. In other words, the goal is to nd topics that are frequent locally in a speci c time frame but those topics, in order to be considered, must be frequent locally in a su cient number of time frames. In this report are described in detail the solution to the just mentioned problem and the relative implementation. Finally, in Section 7, is presented an experimental evaluation of this project using the dataset described in Section 3 and 6 [3].

2 RELATED WORK

In order to design and implement this project, several methods and techniques have been identi ed and implemented; anyway, there are some of them that have been implemented and widely tested but then discarded since they have not produced acceptable results to a possible nal user during the experimental evaluation phase. In the rst part of this section are described all those techniques used in order to pre-process the dataset, while in the second part are described all those methods used to nd consistent topics over time using the pre-processed dataset.

During the pre-processing phase of the dataset have been adopted several techniques used in natural language processing scenarios. The methods that have been implemented and kept in the nal preprocessor script are described in details below and applied following this ordering:

- **Sentence segmentation**: the text string of each tweet is divided into sentences. In other words, the result of this step is a list composed by all the sentences written inside the tweet;
- **Tokenization**: with this technique, each sentence obtained from the previous step is split into even smaller parts called *tokens*. In particular, has been used the blank character as string separator: due to this reason, a token corresponds to a word of the sentence;
- **POS Tagging**: this is the most complex task during the pre-processing of the dataset. The idea of the POS (*Part Of Speech*) technique is to assign to each token the corresponding tag (noun, verbs, adjectives, conjunction, etc.). The collection of tags used (tag set) can be considered as standard since it is used widely in the natural language processing community. In order to implement this technique, the map provided by the Python nltk library [13] has been used.

The other technique used during the pre-processing of the dataset is an entity resolution method that is inspired to *lemmatisation* [7]. The just mentioned technique aims to reduce the in ected form of a word in its canonical form in order to group together di erent words that have the same canonical form. The process implemented in this project is a simpler and non-automatic version of lemmatisation but still very e ective on the given dataset. The goal is to de ne a set of **aliases** stored on the disk in a CSV le interpreted as a dictionary: if a word has an entry in that data structure, it will be replaced with its alias that represents a more general word (see Section 6 for further details).

Another technique that was initially implemented is **stemming** [7]. The aim of this method is to reduce the in ected form of a word in its root form (or "word stem") in order to group together words that have the same root form. The idea was to substitute the given word with its stem during the pre-processing phase. The major problem arose when the script to nd the consistent topics over time has been run: the obtained results were not so meaningful since the root form of a word cannot be presented as a topic to the nal user. Due to this reason and after many

attempt to improve this approach, the just described technique has been abandoned.

In order to find out all the consistent topics over time, the major technique that has been used is the **frequent itemsets** mining method. The goal of this technique is to identify all the sets of items (in the context of this project, set of words) that appears together a sufficient number of times. In order to achieve the goal of this project, the association rules are not computed since they do not produce any valuable information. In a first attempt, the **A-Priori** algorithm [1] was used and implemented. The goal of this procedure is to find out initially frequent individual terms in the dataset and then extending the scope to larger item sets as long as those item sets appear sufficiently often in the database. The A-Priori algorithm leverages the contra-positive *monotonicity* principle: if a word s does not appear in N sets, then no pair that includes s can appear in N sets.

A further attempt was to implement the PCY algorithm or its multihash variation [10], but a more efficient technique that can be parallelized between many machines has been found: **FP-Growth** [6]. The first step of this procedure, similarly to the A-Priori one, is to compute item frequencies and identify frequent items. Then, the second step of FP-Growth is to use a suffix tree structure, called FP-tree, to encode transactions without generating candidate sets explicitly, due to the fact that they are usually expensive to generate. When this second phase terminates the execution, the computed frequent itemsets can be extracted from the FP-tree. With the aim of speed-up the computation distributing the workload among different machines, a parallelizable version of FP-Growth has been used: **Parallel FP-Growth** [8]. This algorithm distributes the work of growing FP-trees based on the suffixes of transactions. Due to this improvement, it is more scalable than a single-machine computation.

The last technique used with the aim to improve the final result presented to the user is a variation of the **closed itemsets** method. The just mentioned technique must be executed after having identified all the frequent itemsets (when the PFP-Growth algorithm has terminated its execution) and its goal is to compact the final output. The idea of closed itemsets is that if there are two frequent itemsets where one is a subset of the other and they have the same frequency, the subset must be discarded. The method implemented in this project follows the same principle, but given two itemsets where one is a subset of the other, the subset is discarded only if it appears in as many time frames as the super set. If they do not appear in the same number of time frames, it means that probably they are different topics so they must be both kept in the result set. The frequency of appearance of the itemsets is not used as a discriminant parameter in order to decide whether to keep or not a subset of items due to the fact that, in any case, a topic composed by many items is much more meaningful than a topic with less items to the final user.

3   PROBLEM STATEMENT

In order to achieve the goal described in Section 1 (i.e., identify consistent topics over time inside tweets), a public available dataset that groups more than 300.000 tweets that contain the hashtag *#covid19* [3] has been used. As many other datasets composed of unstructured data dumped from a social network such as Twitter, this dataset has lots of fields that characterize each tweet (e.g., publisher's username, location and account information, text of the message, etc.).

To identify consistent topics over time using those tweets, the only fields needed are the publication date of each of them and the relative text. In other words, in the working environment a tweet is defined as a tuple composed by the following fields:

- **date**: the date of the tweet publication, expressed as an integer value that represents the number of seconds that have elapsed since the midnight of 1^{st} January 1970;
- **text**: the text of the tweet represented as a list of words. This list is obtained splitting the source text using the blank character " " as separator and processed as described in detail in Section 6.

Therefore, in order to obtain the formal model of this problem, the following sets are been defined:

TWEET = the set of tweets, defined as TIMESTAMP×TEXT where TIMESTAMP $\subset \mathbb{N}_{\geq 0}$ and TEXT = $\{x | x \in \Sigma^*\}$ where Σ is the reference alphabet;

TIMESPAN = the set of all the possible time spans in which the input dataset can be split, defined as a subset of $\mathbb{N}$;

TS-THRESHOLD = the set of all the possible thresholds to identify frequent terms and topics in a single time span, defined as a subset of $\mathbb{N}_{\geq 0}$

GB-THRESHOLD = the set of all the possible thresholds to identify consistent topics over all the time spans, defined as a subset of $\mathbb{N}$

The aim is to model an utility function f defined as

$$f: \text{INPUT} \mapsto \text{TOPIC}$$

where INPUT is defined as

$$\text{TWEET} \times \text{TIMESPAN} \times \text{TSTHRESHOLD} \times \text{GBTHRESHOLD}$$

and TOPIC = $\{x | x \in \Sigma^*\}$ where Σ is the reference alphabet. In other words, TOPIC is the set of all the possible consistent topics over all the identified time spans.

4   SOLUTION

In order to design and implement an efficient solution to the problem shown in Section 3, all the techniques described in the second part of Section 2 have been adopted. Those methods, however, has not been implemented using a naive approach (i.e., the frequent itemsets algorithm has not been executed over the entire dataset and then its results have been processed to reach the final result set).

Since the goal is to find out consistent topics over time, the first piece of the solution is to identify all the necessary time spans in which the initial dataset has to be split. This has been achieved adopting a method that aims to be as flexible as possible not only with the provided dataset [3], but also with any other one pre-processed with the procedure detailed in Section 6. The developed algorithm takes as input two hyper-parameters: *timespan* and *timeunit*. The "timespan" parameter is an integer value greater than zero that identify how long would be considered a single time frame, while the "timeunit" parameter is a string that defined the unit of measure for which the "timespan" parameter has to be evaluated. In the actual solution, the possible time units supported are "day" and "hour"; anyway, other time units can be added quite easily modifying the script module that interpret those parameters and split the dataset entries. As an example, if the time span is set to 1 and the time unit is set to "day", the maximum difference between the timestamps of the last and the first tweet of each time frame would be less or equal than 86400 (i.e., the number of seconds in a single day). As a result of this first phase, it is obtained a list of time frames where each entry

is itself a list of all the tweets included in that time frame. The just mentioned procedure is shown in Algorithm 1.

Then, for each time frame that has been identi ed, the frequent itemsets algorithm is applied. As a rst attempt, was implemented the base A-Priori algorithm, optimizing the overall computation leveraging all the possible CPU cores splitting the workload among di erent processes (the number of processes to start would have been con gurable hyper-parameter of the algorithm and initially, as default value, would have been set to four). This approach seems to be, even from the rst stage of the project development, not a scalable solution since the overall computation has to be performed by a single machine. The problem was in turn exacerbated by the fact that the computation takes relatively long time to compute only the frequent pairs: computing all the super sets of frequent itemsets would have taken lot of computation time even in a multi-process environment. Due to these reasons, this initial solution was implemented and then discarded but it has been taken as "baseline method" during the evaluation phase of the delivered solution as described in Section 7.

The method applied in the delivered solution is the FP-Growth algorithm [6], brie y introduced previously in Section 2. This procedure is widely supported in all the frameworks for massive data mining due to the fact that it is most e cient algorithm that can be executed in a cluster of machines to speed up the computation. More precisely, it has been used a parallelizable version of FP-Growth called Parallel FP-Growth, or simply PFP-Growth, as described in Section 2. This algorithm, for each time frame, takes in input the following three parameters:

- *itemsCol*: the name of the column where the sets of items are located in the provided list. In the actual implementation, the sets of items are located in the "tweets" column;
- *minSupport*: the minimum support for an itemset to be identi ed as frequent. In the actual implementation, this parameter is set to a constant value of 0.01 (i.e., 1%). Then, the frequent itemsets for each time frame will be ltered later as described below;
- *minCon dence*: the minimum con dence for generating the association rules. Since for this task those rules are not needed, the parameter is set to a constant value of 0.6.

As soon as the PFP-Growth procedure has terminated its execution, the frequent itemsets are extracted from the result object and then ltered. With this post-processing ltering, the aim is to remove all the single-item frequent sets and all those frequent itemsets whose frequency is less than a support threshold (TS-threshold) decided by the user as described in Section 3.

When the frequent itemsets have been ltered for a single time frame, is computed the maximal frequency value (i.e., how many times appears the most frequent itemset inside that time frame) and a global hash map that keeps track of all the frequent itemsets obtained among all the various iterations is updated. In other words, each frequent itemset is mapped into a list whose length is equal to the number of time frames and, initially, it is lled with "None" values. When the frequent itemset *i* appears in the time frame *s*, the list assigned to the itemset *i* is updated in position *s*. In particular, in position *s* is assigned a oat value that represents the weight of that topic on a scale between zero and one and then it is normalized in a scale between one and ve. In other words, the importance of a consistent topic is computed as the normalized ratio between the frequency of the itemset *i* and the maximal frequency value computed previously: if the

importance is near one then the topic *i* will not be considered as important, while if the importance is near ve then the topic *i* will be considered very important. The procedure just described is shown in detail in Algorithm 2.

As soon as the global hash map, hereafter *frequencies_dict* has been lled with all the results of every time frame, it has been leveraged in order to compute the consistent itemsets, hereafter topics, over time. As described in Section 3, the algorithm has to be provided with a second threshold value called "global threshold" (GB-threshold) that tells how many time frames a topic has to appear in order to be considered consistent over time. In order to achieve the goal, a new hash map is used called *ranking_dict* where its keys are integer value and the relative value is a list: the key *i* maps to a list that contains all the consistent topics that have appeared in *i* time frames. All the frequent topics that have not appeared in a su cient number of time frames (de ned by the global threshold value), will be discarded.

For each entry in the frequencies_dict hash map, the list of weights is summarized as a single oat value that represents the average weight. Then is created a tuple composed by the frequent topic and the relative average weight. Finally, that tuple is appended to the list associated to the number of time frames in which that topic appears in the ranking_dict hash map as described above.

When the ranking_dict hash map has been built completely, it will start the nal procedure for data manipulation. This latest part has been developed following a precise *implementation choice*: a consistent topic that appears in few time frames is more important than a consistent topic that appears in many time frames. This choice has been taken since those consistent topics that appear frequently in many time frames would be detected by any simple frequent itemsets algorithm, while the purpose of this project is di erent and the relevance of a topic not frequent in the entire dataset but very recurrent in few time frames is much more higher.

After that the above-mentioned implementation choice has been stated, the goal of this nal procedure is to sort by frequency average weight the obtained entries of the ranking_dict hash map. In order to follow the just described property, the ordering procedure starts from those consistent topics that appears in few time frames and then continues with those topics that appear in more time frames. This procedure will terminate its execution when all the frequent topics have been processed or when the maximum number of topics to show to the nal user has been reached. Indeed, the user can specify with a hyper-parameter what is the maximum number of consistent topics over time that have to be shown when the overall procedure terminates. This feature aims to speed up the computation in all those cases where the user has used a low TS-threshold value (i.e., there will be lots of frequent topics) since the just described procedure involves the ordering.

Furthermore, in the procedure just mentioned is applied the variation of the closed itemset technique as described in Section 2. The aim is to remove all those subsets that appear in as many time frames as a super set and compact the output presented to the nal user. This choice has been taken due to the fact that a topic composed by many words is much more meaningful with respect to a topic composed by less words. The just described procedure is shown in detail in Algorithm 3.

As a nal step, the delivered procedure has to build an output le in order to present the obtained results to the nal user. In order to achieve this goal, it has been used the matplotlib Python

---

Algorithm 1 Split dataset by time frame

1: **function** _ _ _ $(dataset, timespan, timeunit)$
2:     $base\_timeframe\_item \leftarrow None$
3:     $previous\_timestamp \leftarrow None$
4:     $start \leftarrow 0$
5:     $end \leftarrow 0$
6:     $timeframes \leftarrow []$
7:     **for** row in dataset **do**
8:         $actual\_timestamp \leftarrow int(row[0])$
9:         **if** base_timeframe_item is None **then**
10:           $base\_timeframe\_item \leftarrow actual\_timestamp$
11:           $previous\_timestamp \leftarrow actual\_timestamp$
12:         **else if** same_timeframe(base_timeframe_item, actual_timestamp) **then**
13:           $previous\_timestamp \leftarrow actual\_timestamp$
14:           $end \leftarrow end + 1$
15:         **else**
16:           $timeframes.append((start, end))$
17:           $start \leftarrow end + 1$
18:           $end \leftarrow end + 1$
19:           $base\_timeframe\_item \leftarrow actual\_timestamp$
20:           $previous\_timestamp \leftarrow actual\_timestamp$
21:         **end if**
22:     **end for**
23:     **if** start != end **then**
24:         $timeframes.append((start, end))$
25:     **end if**
26:     **return** $timeframes$
27: **end function**

Algorithm 2 Find frequent topics

1: **function** _ _ $(timeframes\_limits, dataset, ts\_threshold)$
2:     $frequent\_itemsets\_dict \leftarrow \{\}$
3:     $tf\_id \leftarrow 0$
4:     **for** start, end in timeframes_limits **do**
5:         $tweets\_in\_timeframe \leftarrow dataset.subset(start, end + 1)$
6:         $FPG\_results \leftarrow FP - Growth(tweets\_in\_timeframe, minSupport = 0.01, minConfidence = 0.6)$
7:         *Remove from FPG_results itemsets with length 0 and itemsets with frequency less than ts_threshold*
8:         $max\_freq \leftarrow itemset\_most\_frequent(FPG\_results)$
9:         **for** item in FPG_results **do**
10:           $key \leftarrow item.itemset$
11:           **if** frequent_itemsets_dict not contains key **then**
12:             $frequent\_itemsets\_dict[key] \leftarrow$ *[None for i in range(num_timeframes)]*
13:           **end if**
14:           $values \leftarrow frequent\_itemsets\_dict[key]$
15:           $values[tf\_id] \leftarrow (4 * item.freq/max\_freq) + 1$
16:           $frequent\_itemsets\_dict[key] \leftarrow values$
17:         **end for**
18:         $tf\_id \leftarrow tf\_id + 1$
19:     **end for**
20:     **return** $frequent\_itemsets\_dict$
21: **end function**

---

library [9] that can be leveraged to produce a single PDF report
le that contains all the graphs, one for each time frame that
has been identi ed. With the aim to respect the implementation
choice, the plots have been stored following that principle: the
rst plot contains all those consistent topics that have appeared
in few time frames, the second plot contains all those consistent
topics that have appeared in at least one more time frame with
respect to the previous graph and so on. Furthermore, if the limit

of items has been reached, the procedure will be stopped but
the most relevant consistent topics are already considered in the
nal result set.

The type of plot is the horizontal histogram where on the Y axis
there are the consistent topics and on the X axis there are the
frequency weights on a scale between 1 (low frequency) and 5
(high frequency) and the topics are sorted from the most frequent
(top) to the less frequent (bottom) in the relative number of time

---

Algorithm 3 Find consistent topics over time

1: **function** _ _ _ _ $(frequencies\_dict, gb\_threshold, max\_topics\_to\_show)$
2: $\quad ranking\_dict \leftarrow \{\}$
3: $\quad max\_key \leftarrow 0$
4: $\quad$**for** key in frequencies_dict **do**
5: $\qquad n\_of\_non\_none \leftarrow$ *sum(x is not None for x in frequencies_dict[key])*
6: $\qquad$**if** n_of_non_none >= gb_threshold **then**
7: $\qquad\quad$**if** ranking_dict not contains n_of_non_none **then**
8: $\qquad\qquad ranking\_dict[n\_of\_non\_none] \leftarrow []$
9: $\qquad\quad$**end if**
10: $\qquad\quad sum\_weights \leftarrow 0$
11: $\qquad\quad$**for** entry in frequencies_dict[key] **do**
12: $\qquad\qquad$**if** entry is not None **then**
13: $\qquad\qquad\quad sum\_weights \leftarrow sum\_weights + entry$
14: $\qquad\qquad$**end if**
15: $\qquad\quad$**end for**
16: $\qquad\quad weight\_mean \leftarrow sum\_weights/n\_of\_non\_none$
17: $\qquad\quad ranking\_dict\_value \leftarrow ranking\_dict[n\_of\_non\_none]$
18: $\qquad\quad$*ranking_dict_value.append((weight_mean, key))*
19: $\qquad\quad ranking\_dict[n\_of\_non\_none] \leftarrow ranking\_dict\_value$
20: $\qquad\quad$**if** n_of_non_none > max_key **then**
21: $\qquad\qquad max\_key \leftarrow n\_of\_non\_none$
22: $\qquad\quad$**end if**
23: $\qquad$**end if**
24: $\quad$**end for**
25: $\quad n\_topics\_already\_considered \leftarrow 0$
26: $\quad actual\_key \leftarrow gb\_threshold$
27: $\quad$**while** actual_key <= max_key and n_topics_already_considered < max_topics_to_show **do**
28: $\qquad$**if** ranking_dict contains actual_key **then**
29: $\qquad\quad values \leftarrow sort(ranking\_dict[actual\_key])$
30: $\qquad\quad new\_values \leftarrow []$
31: $\qquad\quad$**for** topic_tuple in values **do**
32: $\qquad\qquad superset\_found \leftarrow false$
33: $\qquad\qquad$**for** topic_to_compare in values **do**
34: $\qquad\qquad\quad$**if** topic_to_compare.length > topic_tuple.length and is_superset(topic_to_compare.length, topic_tuple) **then**
35: $\qquad\qquad\qquad superset\_found \leftarrow true$
36: $\qquad\qquad\quad$**end if**
37: $\qquad\qquad$**end for**
38: $\qquad\qquad$**if** not superset_found **then**
39: $\qquad\qquad\quad$*new_values.append(topic_tuple)*
40: $\qquad\qquad\quad n\_topics\_already\_considered \leftarrow n\_topics\_already\_considered + 1$
41: $\qquad\qquad\quad$**if** n_topics_already_considered == max_topics_to_show **then**
42: $\qquad\qquad\qquad$**break**
43: $\qquad\qquad\quad$**end if**
44: $\qquad\qquad$**end if**
45: $\qquad\quad$**end for**
46: $\qquad\quad ranking\_dict[actual\_key] \leftarrow new\_values$
47: $\qquad$**end if**
48: $\qquad actual\_key \leftarrow actual\_key + 1$
49: $\quad$**end while**
50: $\quad$**return** $ranking\_dict$
51: **end function**

---

frames. The le name of the output le can be decided by the nal user using a hyper-parameter or, if that parameter is not used, the output le would be "report.pdf".

5 IMPLEMENTATION

In order to implement the delivered project, di erent tools have been used. First of all, the entire project has been developed using the Python 3 programming language due to the fact that it has many libraries freely available that help programmers to deal with huge quantity of data and process them in relatively easy manner using its natural language toolkit frameworks. Thanks to those libraries, it has been possible to develop and test many slight variation of the delivered solution with the aim to nd the most e cient one in a relative short time. As a result, one of the many scripts that have been developed to compare the e ciency

has been taken as a quite advanced "baseline method" during the experimental evaluation phase described in Section 7.

The main external tool used during the development of this project is *Apache Spark* [16] in its version 3.0.1, a uni ed analytics engine for large scale data processing. Apache Spark is an e cient, fast and general purpose cluster computing framework with high-level APIs in di erent programming languages. Since it enables in-memory computations, Spark is way more e cient with respect to other frameworks that have been created for the same purpose, such as Apache Hadoop [5]. The Apache Spark strength is that it provides many APIs that permits to parallelize the workload among all the cluster of machines in a really easy manner. Furthermore, Spark's APIs permits to access its data structures using SQL-like methods or using data-frame operations that are more or less the same as those ones exposed by the Pandas Python library [14].

In order to integrate Apache Spark in the Python 3 script, the *Pyspark* library [12] has been used. Pyspark has been released in order to integrate Apache Spark and the Python programming language: it actually is a Python API for Spark. In particular, in order to support Apache Spark that has been written using the Scala programming language, Pyspark allows the Python interpreter to dynamically interface itself with the JVM (Java Virtual Machine) objects necessary for the script execution. This has been achieved using the Py4J library [11] that is inspired to the Java Remote Procedure Call standard.

To implement the solution described in Section 4, two modules of Pyspark have been used: *PysparkSQL* and *MLlib*. The PysparkSQL module allows to apply SQL-like operations on a huge amount of structured or semi-structured data; indeed SQL queries can be used to manage its data structures. This module can be interfaced with several data sources such as Apache Hadoop, MySQL or Parquet and it can load even external data such as, for example, CSV or JSON les. Furthermore, this module introduces the DataFrame data structure, a tabular representation of structured data that is similar to that of a table from a relational database management system. Since all the datasets that will be used with the delivered solution are CSV les, the choice to use the PysparkSQL module has been quite obliged. In fact, the input dataset is immediately parsed and transformed into a Pyspark data frame object and all the operations of data manipulation are executed on that data structure due to the fact that those operations can be parallelized among di erent machines.

On the other hand, the MLlib module of Pyspark is a wrapper for the Spark's machine learning library and it leverages data parallelism techniques to store and deal with huge amount of data. The MLlib module supports many machine-learning algorithms for classi cation, regression, clustering, collaborative ltering, dimensionality reduction, and underlying optimization primitives. This module has been used due to the fact that it supports directly the PFP-Growth algorithm described in Section 4 and it uses the same hyper-parameters already presented. As a point of strength, the MLlib module manages the parallelization on multiple machines itself, hence the developer do not have to worry about this feature.

The other tool that has played a crucial role during the development of the nal solution is the nltk Python library (Natural Language ToolKit) [13]. This module has been developed in order to deal in an e cient and e ective way with human language data. Indeed, it has been used in order to parse and interpret the words of each tweet inside the given dataset. Furthermore, all the natural language processes described in Section 2 (sentence segmentation, tokenization and POS tagging) have been used thanks to the fact that the nltk library implements directly those procedures and it provides for free a POS tagging map for english words.

The third tool that has been used to build the delivered solution is the matplotlib Python library [9]. This module has been used in order to present in a clear and concise manner the results obtained at the end of the script execution to the nal user. In particular, the matplotlib library has been used to generate a single PDF le that contains several plots where each of them represents a subset of all the consistent topics over time that have been detected as described in details in Section 4.

The latest tool used in both dataset preprocessor script and in the consistent topics nder script is the atpbar Python library [2]. The goal of this module is to show to the user a progress bar with the aim to inform the latter about the state of the execution progress and, in particular, how many tweets there are in the provided dataset and how many time frames have been identi ed. The strength of this library is that it can show to the user progress bars that are managed by di erent processes. Due to this fact, the nal user will be always aware about the execution progress of both the above mentioned scripts even if they have been executed in a multi process environment.

6 DATASET

As described previously in Section 3, in order to build and test the solution presented in Section 4 a public available CSV dataset with more than 300.000 tweets that contain the hashtag *#covid19* [3] has been downloaded from kaggle.com and processed. As many other datasets that group together huge quantity of unstructured data, it had to undergo a pre-processing phase in order to remove unnecessary features and noisy data.

In that dataset, a tweet is de ned as a data item composed of thirteen features: *user-name*, *user-location*, *user-description*, *user-created*, *user-followers*, *users-friends*, *user-favourites*, *user-veri ed*, *date*, *text*, *hashtags*, *source*, *is-retweet*. In order to achieve the goal described in Section 3, only two features are needed: the publication date of the tweet and the relative text. Due to this reason, the rst stage of the pre-processing consists to cut o all the non-relevant data features. This has been achieved quite easily since the dataset is CSV (Comma Separated Value) le and the Pandas Python library [14] has lots of APIs to manage e ciently such big datasets.

After this initial stage, each tweet has been processed singularly in a multi-process execution where to each process is assigned a portion of the input dataset. For each tweet in the dataset portion, each process has to manipulate both the date and the text elds in the following manner:

Date eld is transformed into an UNIX timestamp. In particular the publication date is stored in the input dataset as a string formatted as "yyyy-mm-dd hh:mm:ss", but thanks to the *datetime* Python library this transformation can be performed in a easy and fast way. With the aim to be complaint also with the format of another dataset that groups together tweets in order to perform sentiment analysis [15], the script responsible for the date manipulation is capable to transform also dates formatted like "Mon Apr 06 22:19:45 PDT 2009" into a UNIX timestamp.

Text eld is transformed into a list of words. In particular, thanks to the *nltk* [13] Python library, the initial text has been split using the blank character " " as separator and

only the useful words that can be leveraged in order to derive a topic have been kept. Due to this reason, only nouns and adjectives appear in the final result list. Furthermore, there are other situations where a word, even it is a name or an adjective, cannot be considered:

- if the word contains a slash "/", non-ASCII characters, special Unicode sequences;
- if the word represents a numeric value;
- if the word is composed by only one character;
- if the word is a stop word. This method has been implemented because there are some words like "https" or "amp" that are tagged as nouns but they cannot be leveraged to build a topic.

As a note, if the preprocessor script has been executed in the debug mode, all the non-considerable words are dumped inside a separated CSV file (together with the relative timestamp) in order to check if the script cut off useful words. Then, each considerable word is filtered and sanitized in order to remove all the eventual noisy characters. In particular, the filtering method performs the following operations:

(1) all the word's characters are lowered;
(2) all the emojis are removed;
(3) all the non-alphanumeric characters are substituted with a blank space (e.g., the word "white-house" is transformed to "white house");
(4) the word is split again using the blank character as separator to identify eventual other words after the previous operation;
(5) all the obtained words are rechecked again to see if they are considerable and, eventually, discarded;
(6) as a final operation, each of the obtained words are checked with an aliases map. If there is an entry for a word, then the associated value is substituted. The aim is to generalize as much as possible the words used inside the tweets leveraging the knowledge about that dataset in order to obtain better results: since all the texts contain the word "covid19" or one of its many variations (e.g., "sars-cov-2", "covid", "coronavirus", etc.), all of them are mapped to the word "covid19".

After these filtering operations, a final check is performed over all the considerable words in order to find out if there are doubled words or empty string.

When each process has terminated its execution, all the partial results are collected in a single list by the master process using a pipe-based communication and sort that list in ascending order using the timestamp field as key. The result is a list of tweets ordered by their publication: this feature can be leveraged by the main script described in Section 5 in order to lower the overall complexity. If the preprocessor script has been executed in debug mode, also the list that contains all the discarded words would be collected but not sorted.

As a final step, the result list has been converted by the master process in a Pandas [14] dataframe object and stored in a separated CSV file. The same procedure would be applied to the second list with all the discarded words if the preprocessor script has been executed in debug mode.

7 EXPERIMENTAL EVALUATION

7.1 Test dataset

In order to test the speed and the quality of the output produced by the delivered script, the dataset that groups more than 300.000 tweets that contain the hashtag *#covid19* [3] has been used massively. This dataset has been already shown in details in Section 6.

With the aim of perform a more in-depth analysis of the performances of the script, a bigger dataset has been uses: this dataset [15] contains 1.600.000 tweets extracted using the twitter API and they have been annotated in order to be used to detect sentiment. This dataset is stored in a CSV file where each item is composed by six fields: *target* (the polarity of the tweet), *ids* (tweet's id), *date*, *flag*, *user* and *text*.

The most significant problem with using this dataset was during the pre-processing phase: the date field of each tweet was formatted using a not too handy format. As an example, the date of the first tweet was "Mon Apr 06 22:19:49 PDT 2009" and it has been quite hard to transform that date to a UNIX timestamp as required.

The other problem about this dataset is that it has not too many consistent topics over time, even if the quality of those detected by the algorithm described in Section 4 is high. As a matter of fact, the hyper-parameters TS-threshold and GB-threshold has been set to, respectively, 5 and 2 in order to catch some consistent topics (with the dataset about covid19 [3] those hyper-parameters have been set usually to, respectively, 70 and 3).

REFERENCES

[1] apriori-algorithm [n.d.]. *A-Priori algorithm reference.* Retrieved January 10, 2021 from https://en.wikipedia.org/wiki/Apriori_algorithm

[2] atpbar-reference [n.d.]. *Atpbar library reference.* Retrieved January 10, 2021 from https://pypi.org/project/atpbar/

[3] Covid19-Dataset [n.d.]. *Tweets with the hashtag covid19.* Retrieved December 17, 2020 from https://www.kaggle.com/gpreda/covid19-tweets

[4] Customer-Complaints [n.d.]. *Six Steps to Dealing with Customer Complaints.* Retrieved December 18, 2020 from https://www.eonetwork.org/octane-magazine/special-features/sixstepstodealingwithcustomercomplaints

[5] hadoop-reference [n.d.]. *Apache Hadoop reference.* Retrieved January 10, 2021 from https://hadoop.apache.org/

[6] Jiawei Han, Jian Pei, and Yiwen Yin. 2000. Mining Frequent Patterns without Candidate Generation. *SIGMOD Rec.* 29, 2 (May 2000), 1–12. https://doi.org/10.1145/335191.335372

[7] lemmatisation-stemming [n.d.]. *Lemmatisation and stemming reference.* Retrieved January 10, 2021 from https://nlp.stanford.edu/IR-book/html/htmledition/stemming-and-lemmatization-1.html

[8] Haoyuan Li, Yi Wang, Dong Zhang, Ming Zhang, and Edward Y. Chang. 2008. Pfp: Parallel Fp-Growth for Query Recommendation. In *Proceedings of the 2008 ACM Conference on Recommender Systems (RecSys '08)*. Association for Computing Machinery, New York, NY, USA, 107–114. https://doi.org/10.1145/1454008.1454027

[9] matplotlib-python [n.d.]. *Matplotlib Python library reference.* Retrieved January 10, 2021 from https://matplotlib.org/

[10] pcy-multihash-algorithms [n.d.]. *PCY and multihash algorithms reference.* Retrieved January 10, 2021 from https://medium.com/weekly-data-science/the-pcy-algorithm-and-its-friends-ecba67216190

[11] py4j-reference [n.d.]. *Py4J library reference.* Retrieved January 10, 2021 from https://www.py4j.org/

[12] pyspark-reference [n.d.]. *Pyspark Python library reference.* Retrieved January 10, 2021 from https://spark.apache.org/docs/latest/api/python/index.html

[13] Python-nltk [n.d.]. *Natural Language ToolKit python library.* Retrieved January 10, 2021 from https://www.nltk.org/

[14] Python-Pandas [n.d.]. *Pandas python library.* Retrieved January 10, 2021 from https://pandas.pydata.org/

[15] Sentiment-Analysis-Dataset [n.d.]. *Tweets for sentiment analysis.* Retrieved January 10, 2021 from https://www.kaggle.com/kazanova/sentiment140

[16] spark-reference [n.d.]. *Apache Spark reference.* Retrieved January 10, 2021 from https://spark.apache.org/

[17] Twitter-statistics [n.d.]. *Twitter usage statistics.* Retrieved January 10, 2021 from https://www.internetlivestats.com/twitter-statistics/

[18] upGrad [n.d.]. *Benefits and Advantages of Big Data and Analytics in Business.* Retrieved December 17, 2020 from https://www.upgrad.com/blog/

benefts-and-advantages-of-big-data-analytics-in-business/